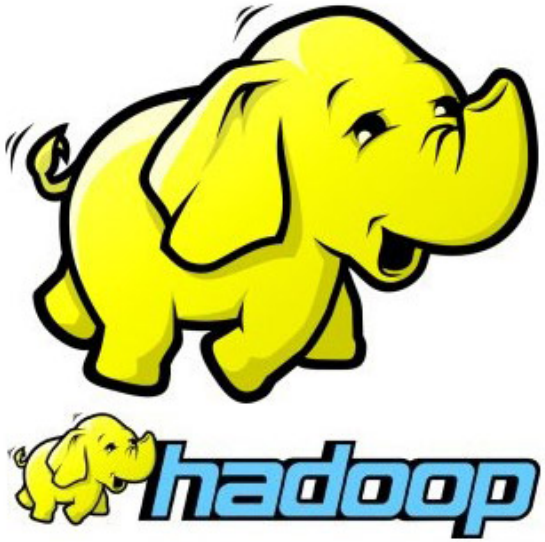

Yaroslav Yermilov

Disclaimer - no words on slides

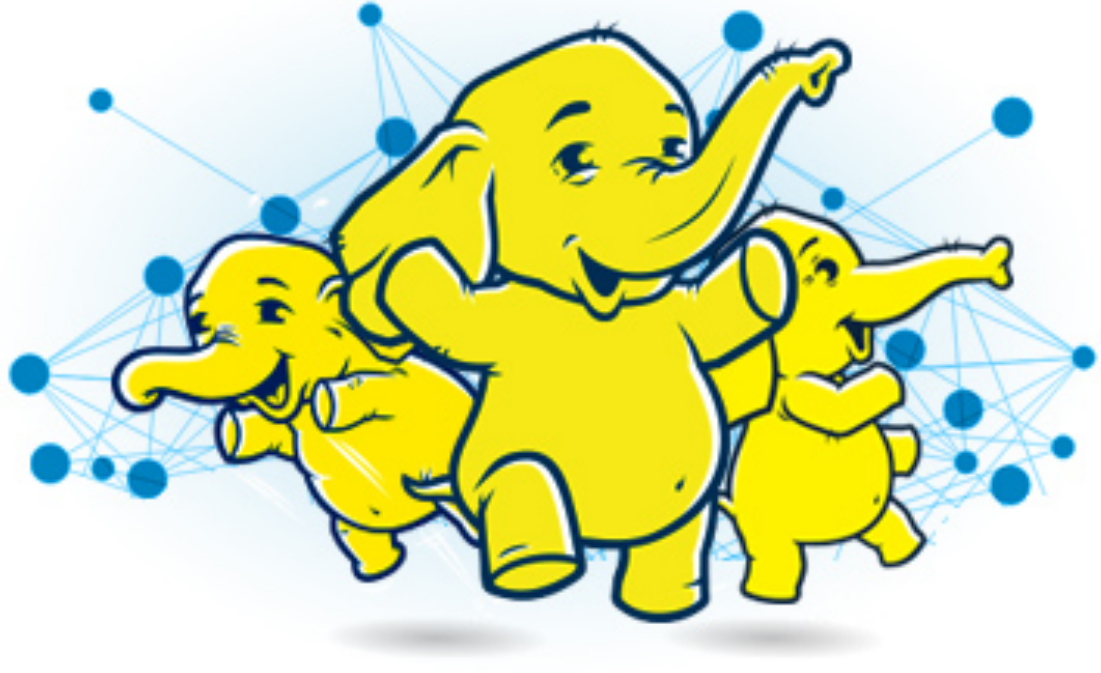

Hadoop MapReduce

MapReduce model

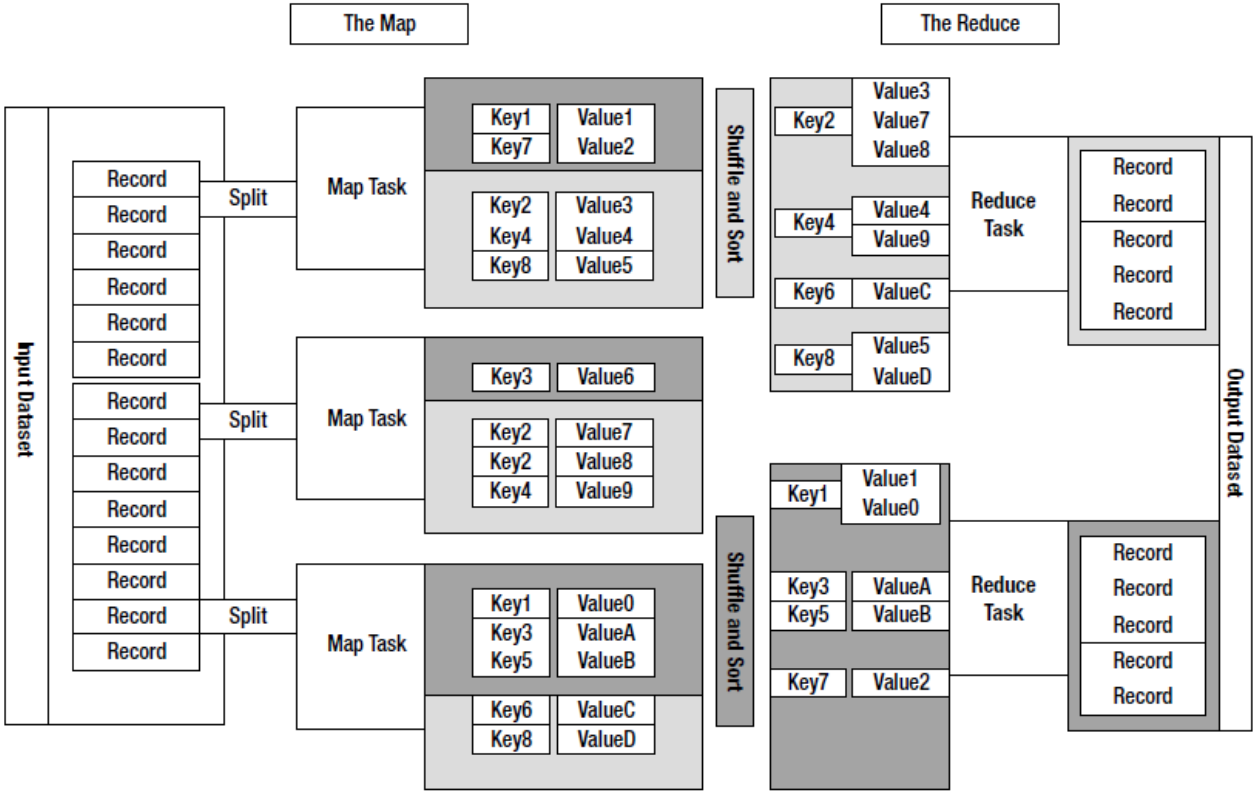

Web crawler example

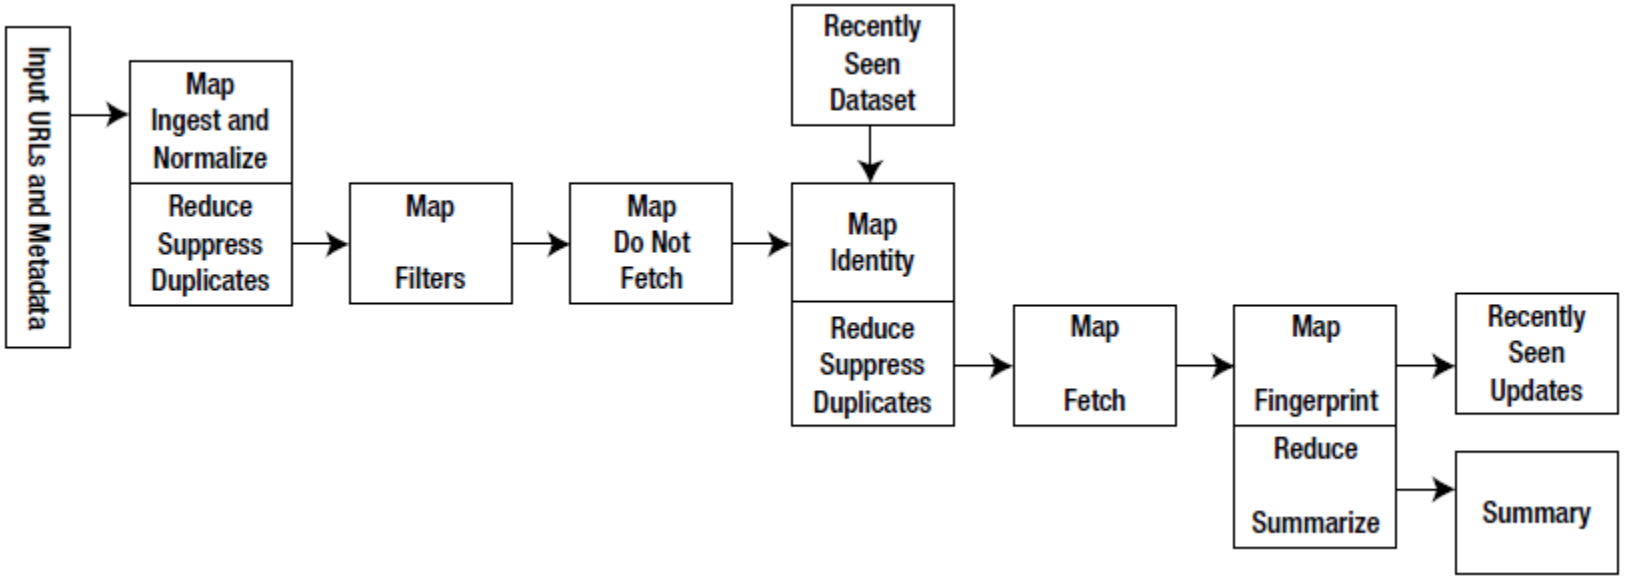

Parts of MapReduce job

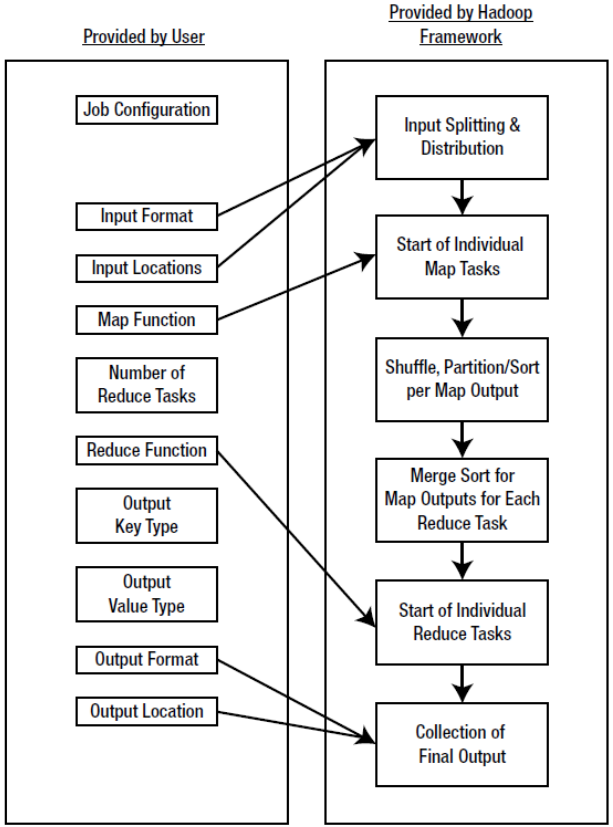

Hadoop cluster

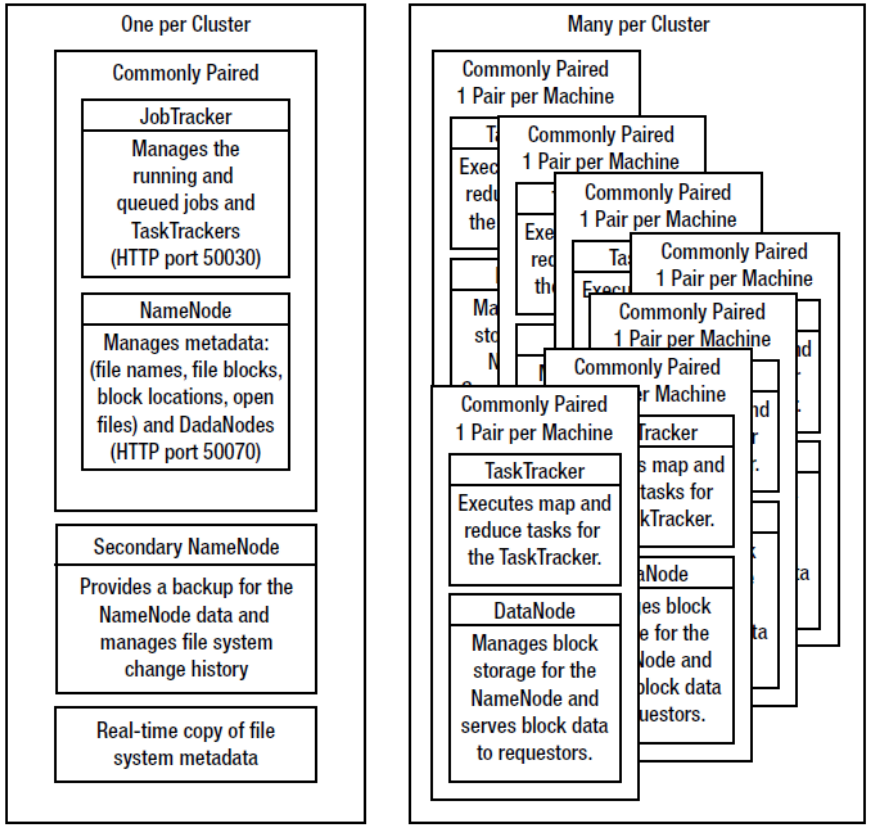

How mapper works

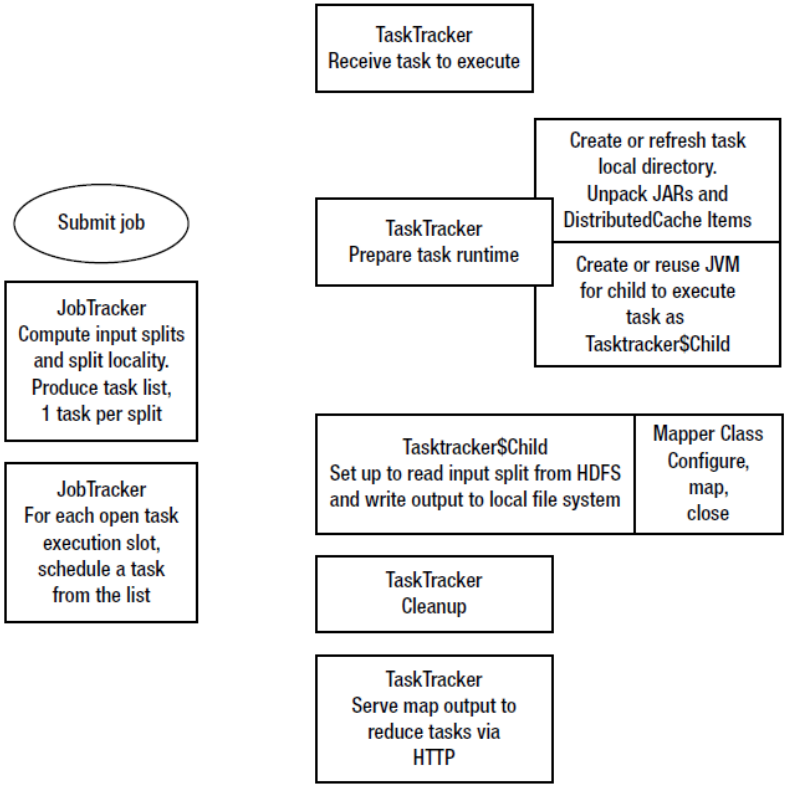

How reducer works

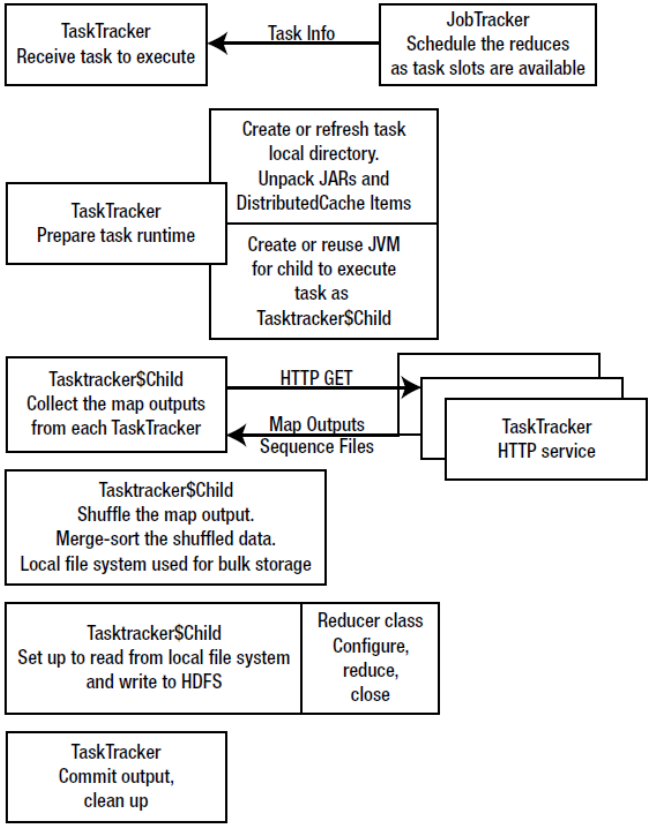

Hadoop HDFS

How HDFS works

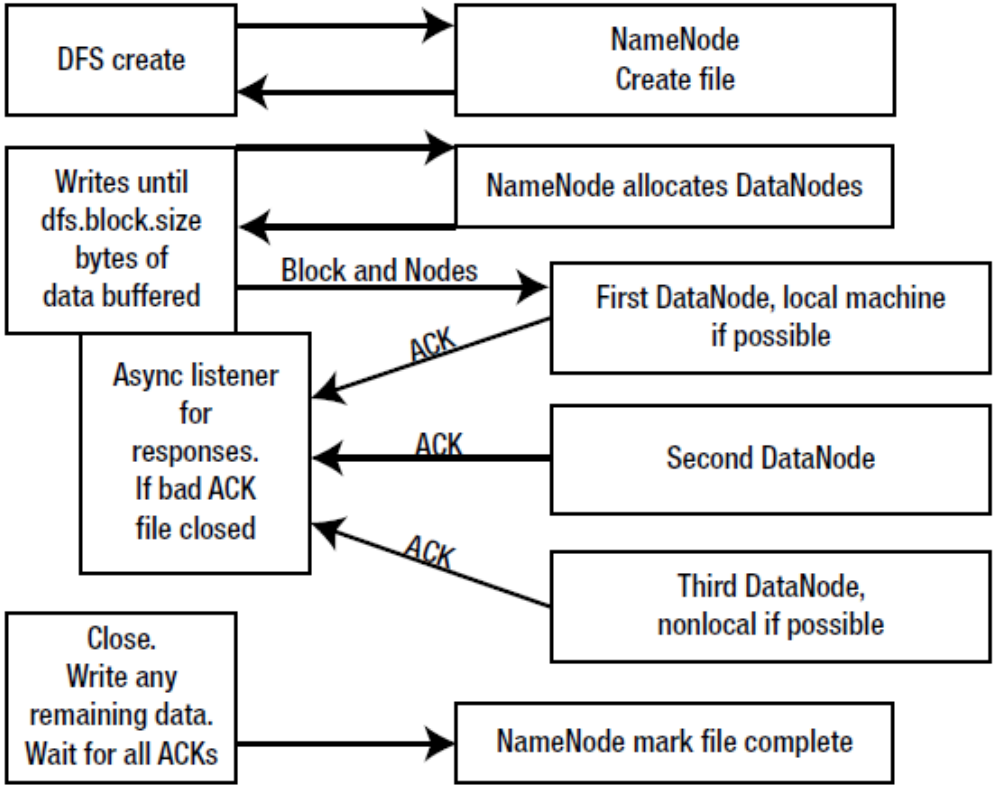

Hive


```
SELECT a.year, a.player_id, a.runs FROM batting a
JOIN (SELECT year, max(runs) runs FROM batting GROUP BY year ) b
ON (a.year = b.year AND a.runs = b.runs) ;
```

Pig

```
batting = load 'Batting.csv' using PigStorage(',');
runs = FOREACH batting GENERATE $0 as playerID, $1 as year, $8 as runs;
grp_data = GROUP runs BY (year);
max_runs = FOREACH grp_data GENERATE group as grp, MAX(runs.runs) as max_runs;
join_max_run = JOIN max_runs BY ($0, max_runs), runs by (year,runs);
join_data = FOREACH join_max_run GENERATE $0 as year, $2 as playerID, $1 as runs;
dump join_data;
```

HBase

Sqoop

Flume

Zookeeper

Oozie

Falcon

Mahout

Avro

Sources

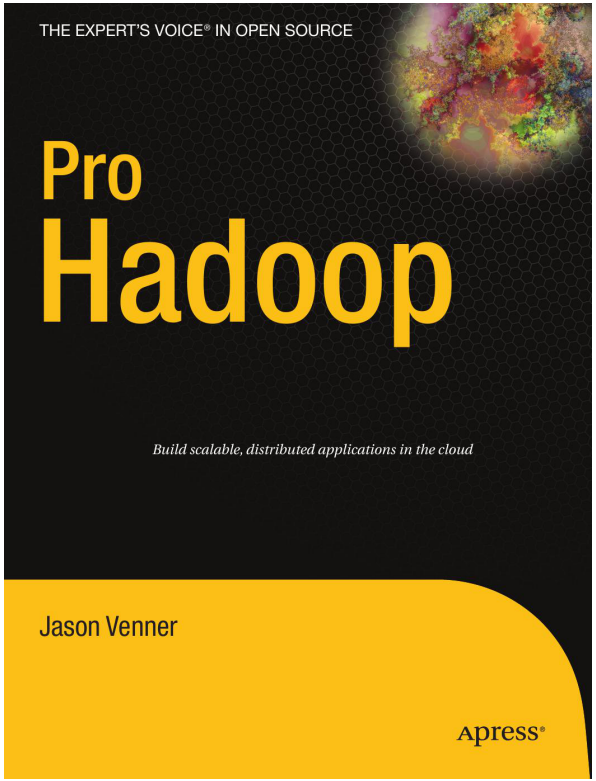
THE EXPERT'S VOICE® IN OPEN SOURCE
Pro
Hadoop
Build scalable, distributed applications in the cloud
Jason Venner
Apress®